Normalizing flows: from change of variables to learned densities

Teacher's guide and worked solutions

Bayesian machine learning | Skoltech

A normalizing flow transforms a simple random variable through an invertible map. Its density follows from the change-of-variables formula. The seminar builds this calculation from a scalar affine transformation, then uses affine coupling layers and their compositions to learn a two-dimensional density. Three implementation tasks lead to a trained RealNVP model and a held-out comparison with a fitted Gaussian. The optional Bayesian appendix uses the same model family as a variational approximation to a curved, posterior-shaped analytic benchmark. This benchmark is specified directly, rather than obtained as the posterior of a fitted data model.

The main distinction to preserve throughout is between *evaluating the model density exactly* and *representing the desired distribution exactly*. The Jacobian formula gives the density of the chosen transformation, up to ordinary floating-point error. It does not remove approximation, estimation, or optimization error in learning that transformation.

The companion student and solutions notebooks have the same lesson order. The student notebook contains three unfinished implementations, supplied checks, experiments, and interpretation questions. The solutions notebook contains the worked code and executed figures. This guide adds board derivations, suggested hints, answers to oral questions, and a route through the material in class.

1 Teaching plan

1.1 The 90-minute lesson

Exercise time is included in each block. The main lesson ends with a trained density model; neither optional appendix is required to reach that endpoint.

Clock	Topic	Student activity and intended result
0–15	Random variables, maps, and densities	Derive the volume correction; fix the sampling and encoding directions.
15–25	Task 1: scalar affine flow	Implement the transformation and density; check the sign and the inverse.
25–40	Affine coupling	Derive the inverse and triangular Jacobian; explain why the neural network need not be invertible.
40–55	Task 2: coupling layer	Implement both directions and per-sample log determinants; inspect identity initialization and checks.
55–70	Composition and Task 3	Reverse layer order for encoding; implement `log_prob` and sampling.
70–85	Likelihood training and diagnostics	Run supplied training; compare held-out log densities, samples, densities, and encoded observations.
85–90	Wrap-up	Explain what is exact, what is learned, and how this connects with sampling and the Laplace approximation.

Before the seminar. Run the complete solutions notebook on the classroom machine. The main lesson uses PyTorch, NumPy, Matplotlib, and the notebook display tools; all data are generated locally. Confirm that the training experiment completes on CPU and retain its executed outputs. Use the student notebook for projection and keep the solutions available separately. Students need the chain rule, determinants, Gaussian densities, expectations, and basic tensor operations. Review that a matrix multiplication and an elementwise multiplication are different operations.

During the seminar. Follow the cells in order. Students write the three mathematical implementations; data generation, plotting, optimization, and evaluation helpers are supplied. Ask pairs to switch the typing role between tasks. A function that raises `NotImplementedError` is unfinished; the notebook skips its dependent experiments. Other failures should remain visible, because a shape error or a failed Jacobian check is useful evidence.

Protect the density-estimation endpoint. Start coupling by minute 25, composition by minute 55, and the training comparison by minute 70. If an implementation remains unfinished, show its missing mathematical step and continue with the supplied working solution. Shorten neural-network mechanics and plotting-code discussion before removing the inverse calculation or the held-out comparison. Use saved outputs if runtime is longer on students' machines. Do not add the variational appendix to the timed lesson by default.

1.2 A 60-minute route

Use minutes 0–10 for the change of variables and notation, 10–17 for Task 1, 17–27 for the coupling derivation, 27–37 for Task 2, 37–47 for composition and Task 3, 47–57 for the supplied training results and diagnostics, and 57–60 for the closing questions. Show the automatic-differentiation Jacobian checks without deriving their implementation. Reduce task coding to filling the transformation and log-determinant expressions. Show an already executed training run if necessary. Keep both directions, all three task contracts, and the difference between exact density evaluation and a learned fit.

1.3 Board layout and notation

Keep the convention below visible. Different references call different directions "forward"; identify the actual map before comparing formulas. In this seminar `inverse=False` always means the generative direction.

Symbol or object	Meaning
$z \in \mathbb{R}^d$, p_0	Base variable and its standard-normal density.
$x \in \mathbb{R}^d$, q_ϕ	Observation-space variable and the flow's density.
$x = g_\phi(z)$	Generative map: easy base samples become model samples.
$z = f_\phi(x) = g_\phi^{-1}(x)$	Encoding map used to evaluate the density at an observation.
$\log \lvert \det Dh \rvert$	Log absolute determinant of the map actually applied, not always of the generative map.
Batch arrays	Main examples have shape $(B, 2)$; log densities and coupling log determinants have shape $(B,)$.
$m \in \{0, 1\}^d$	Mask; a one identifies a coordinate preserved by a coupling layer.

The essential board formula is

$$\log q_\phi(x) = \log p_0(f_\phi(x)) + \log \lvert \det Df_\phi(x) \rvert = \log p_0(z) - \log \lvert \det Dg_\phi(z) \rvert, \quad z = f_\phi(x).$$

The two expressions use Jacobians of opposite directions, evaluated at corresponding points. Their signs differ for that reason.

2 From transformed samples to a density

2.1 Conservation of probability

Let Z have density p_0 and let $g : \mathbb{R}^d \to \mathbb{R}^d$ be a differentiable bijection with differentiable inverse f and nonzero Jacobian determinant. Writing $X = g(Z)$ defines a distribution. To calculate its density, take

a small region around z. Locally, the map is approximately linear:

$$g(z+\delta) \approx g(z) + Dg(z)\delta.$$

A linear map changes volume by its absolute determinant. Thus an infinitesimal volume transforms as $dx = |\det Dg(z)|\, dz$. Probability mass is preserved:

$$q(x)\, dx = p_0(z)\, dz.$$

Substitution gives

$$q(x) = \frac{p_0(z)}{|\det Dg(z)|} = p_0(f(x))|\det Df(x)|.$$

The inverse function theorem gives $Df(g(z)) = [Dg(z)]^{-1}$, so the two forms agree. Taking logarithms produces the board formula above. Smooth neural maps satisfy these conditions when constructed as below; with piecewise differentiable networks, the corresponding Jacobian formula applies almost everywhere.

The factor is an absolute determinant because volumes cannot be negative. In one dimension, a decreasing transformation still has a valid positive density. Our affine scales are positive exponentials, so their derivatives are positive; this does not remove the absolute value from the general theorem.

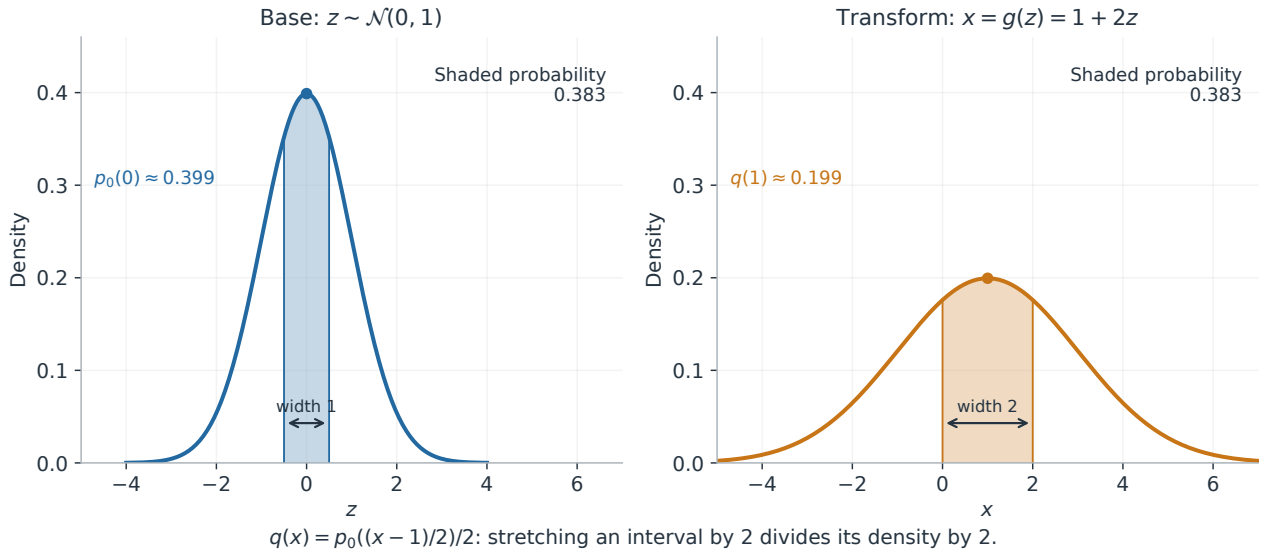


Figure 1: An invertible map transports probability mass and changes volume. Expanding a region lowers its density per unit volume; contracting it raises that density. The Jacobian determinant supplies the correction.

Oral prompt. If a transformation doubles the length of every small interval, what happens to the density? It is divided by two at the corresponding image points. Merely replacing z by $f(x)$ in the base density does not account for the change of units.

2.2 Why the density is normalized

Invertibility and a normalized base imply

$$\int_{\mathbb{R}^d} q(x)\, dx = \int_{\mathbb{R}^d} p_0(f(x))|\det Df(x)|\, dx = \int_{\mathbb{R}^d} p_0(z)\, dz = 1.$$

No partition function is estimated after training. This is a structural property of the model family. It is separate from whether q_ϕ resembles the distribution that generated the observations.

Ordinary sampling from a fitted flow requires independent base draws and one application of g_ϕ to each draw. Conditional on the fitted parameters, these are independent draws from q_ϕ. They are draws from the desired target only to the extent that q_ϕ equals that target. This separates a flow sampler from an MCMC procedure designed to preserve a specified posterior.

3 Task 1: a scalar affine flow

3.1 What students should do

Implement the two functions

```
affine_1d(value, shift, log_scale, inverse=False)
affine_log_prob(x, shift, log_scale)
```

Both directions of the transformation return the transformed tensor and the log absolute derivative of the direction actually applied. For this scalar exercise, that second output has the same shape as the input; it is not summed over a batch. The shift and log scale may be Python scalars, so convert them to tensors with the input's dtype and device before using tensor operations.

3.2 Derivation and a numerical example

Write μ for the shift and a for the log scale:

$$x = g(z) = \mu + e^a z, \qquad z = f(x) = (x - \mu)e^{-a}.$$

The forward derivative is e^a and the inverse derivative is e^{-a}. Thus

$$\log |g'(z)| = a, \qquad \log |f'(x)| = -a.$$

For $Z \sim \mathcal{N}(0, 1)$,

$$\log q(x) = -\tfrac{1}{2}z^2 - \tfrac{1}{2}\log(2\pi) - a, \qquad z = (x - \mu)e^{-a}.$$

This is $\mathcal{N}(\mu, e^{2a})$: e^a is the standard deviation, whereas e^{2a} is the variance. Parameterizing a positive scale by its logarithm keeps it positive for every finite parameter value.

With $\mu = 1$ and $a = \log 2$, base points $(-1, 0, 1)$ become $(-1, 1, 3)$. The forward log derivative is $\log 2$, the inverse log derivative is $-\log 2$, and the density at the center is $1/(2\sqrt{2\pi})$. Its log density is approximately -1.61209. The inverse recovers the original three points.

3.3 Worked implementation

```
def affine_1d(value, shift, log_scale, inverse=False):
    shift = torch.as_tensor(shift, dtype=value.dtype, device=value.device)
    log_scale = torch.as_tensor(log_scale, dtype=value.dtype,
                                device=value.device)
    if inverse:
        transformed = (value - shift) * torch.exp(-log_scale)
        logdet = -log_scale
    else:
        transformed = shift + torch.exp(log_scale) * value
        logdet = log_scale
    return transformed, torch.ones_like(value) * logdet

def affine_log_prob(x, shift, log_scale):
    z, inverse_logdet = affine_1d(x, shift, log_scale,
                                  inverse=True)
    base_log_prob = -0.5 * (z.square() + math.log(2 * math.pi))
    return base_log_prob + inverse_logdet
```

3.4 Checks, hints, and model answers

The checks establish transformed values at two nonunit scales, forward–inverse reconstruction, opposite log determinants at matching points, and agreement with the analytic Gaussian density. Numerical

integration of the density checks normalization independently of reconstruction. An analytic log-scale gradient is compared with automatic differentiation to verify that tensor conversion and the inverse calculation preserve gradients. In the executed reference run, 12,000 draws have sample mean 1.0021 and sample standard deviation 2.0042, close to their analytic values 1 and 2.

A useful first hint is to write the inverse on paper. The second is to identify which derivative each branch returns. For the density, evaluate the base at the encoded point and *add the inverse* log determinant. Do not subtract it again.

Interpretation answer. Increasing a stretches the distribution and reduces its density at corresponding points by e^{-a}. It changes both the encoded coordinate and the normalization term at a fixed x. Changing μ translates the distribution without changing the Jacobian. The determinant correction is what makes the translated and scaled density integrate to one.

4 Affine coupling: an invertible map from an ordinary network

4.1 A split into preserved and transformed coordinates

Partition coordinates into A and B. A coupling layer preserves z_A and uses a network evaluated at that preserved part to transform z_B:

$$x_A = z_A, \qquad x_B = z_B \odot \exp(s(z_A)) + t(z_A).$$

The operations on B are elementwise. The functions s and t can be ordinary neural networks; they need not be invertible. They produce transformation parameters after the preserved coordinates are known.

The inverse is explicit:

$$z_A = x_A, \qquad z_B = (x_B - t(x_A)) \odot \exp(-s(x_A)).$$

The same network is evaluated at the same preserved values in both directions. We solve the affine equation for the changed coordinates. We never solve for the input of the neural network.

4.2 The Jacobian and its determinant

Order the coordinates as (A, B). The generative Jacobian has block form

$$Dg(z) = \begin{pmatrix} I & 0 \\ * & \operatorname{diag}(e^{s(z_A)}) \end{pmatrix}.$$

The lower-left block includes derivatives of s and t. It can be dense, but it does not enter the determinant of this block-triangular matrix. Therefore

$$\det Dg(z) = \prod_{j \in B} e^{s_j(z_A)}, \qquad \log|\det Dg(z)| = \sum_{j \in B} s_j(z_A).$$

The inverse log determinant is the negative of the same sum. Translational terms affect the map and its off-diagonal derivatives; they do not affect this determinant. Autodifferentiation through the explicit inverse still trains both networks through the full log density.

Board example. In two dimensions, preserve z_1 and write $x_2 = e^{s(z_1)} z_2 + t(z_1)$. Its Jacobian is

$$\begin{pmatrix} 1 & 0 \\ e^{s(z_1)} z_2 s'(z_1) + t'(z_1) & e^{s(z_1)} \end{pmatrix}.$$

The determinant is $e^{s(z_1)}$ even when the lower-left term is large. The layer can bend a cloud through the dependence of the changed coordinate on the preserved one, while retaining an inexpensive determinant.

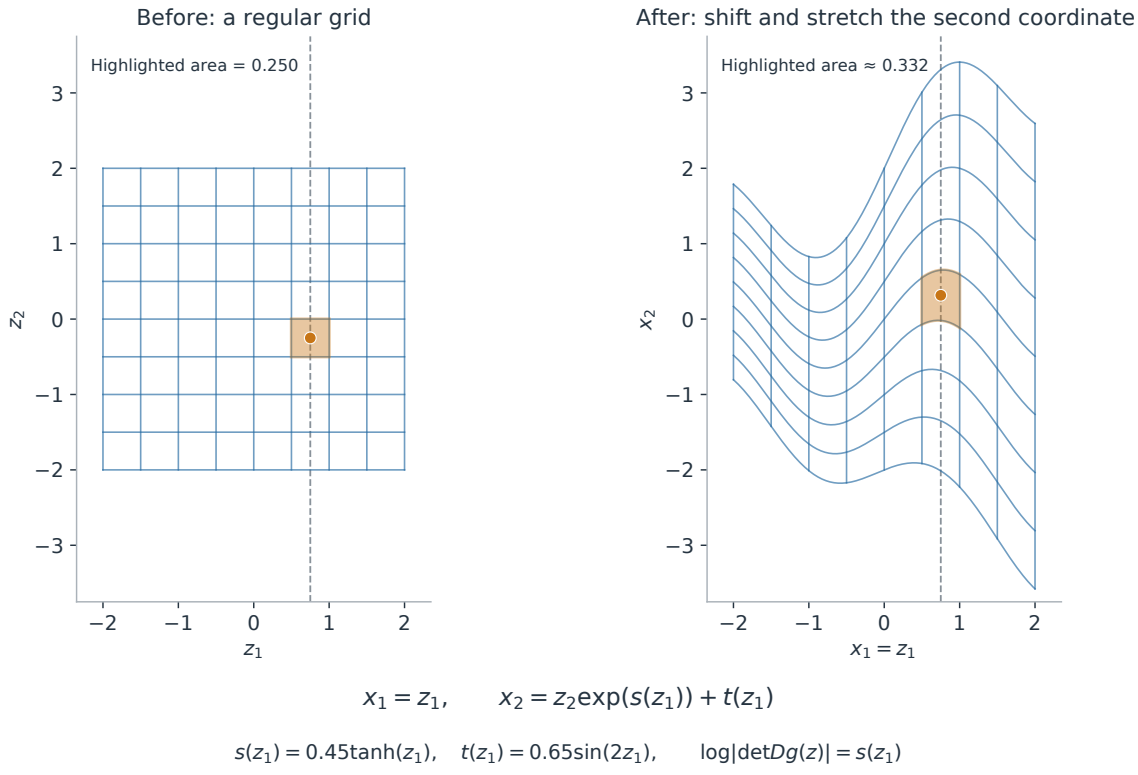


Figure 2: A coupling layer preserves one coordinate and changes the other through a scale and shift determined by the preserved coordinate. A nonlinear shift can bend the sample cloud without requiring inversion of the network.

4.3 Masks and bounded log scales

The implementation keeps full-dimensional tensors and uses a binary mask m. A one means "preserve this coordinate". Let

$$u = m \odot z, \qquad (s_{\text{raw}}, t_{\text{raw}}) = \text{net}(u),$$

where the network returns $2d$ entries split into two vectors of length d. Set

$$s = 1.5 \tanh(s_{\text{raw}}) \odot (1 - m), \qquad t = t_{\text{raw}} \odot (1 - m).$$

Then

$$g(z) = m \odot z + (1 - m) \odot (z \odot e^s + t).$$

The network must receive the masked input. If its scale or shift depends on a coordinate being transformed, this explicit inverse and triangular Jacobian calculation no longer follow.

The bound limits each changed coordinate's log scale to $(-1.5, 1.5)$. Its scale is consequently between approximately 0.223 and 4.482. This controls one source of numerical difficulty; it does not guarantee stable gradients or a good fit for an arbitrary composition. The shift is not bounded by this construction. Register the mask as a PyTorch buffer so it moves with the module between devices and is saved with its state.

5 Task 2: implement a coupling layer

5.1 What students should do

Complete the coupling layer's `transform` method. The argument `inverse=False` selects generation; `inverse=True` selects encoding. The supplied constructor stores the mask and network as `self.mask`

and `self.net`. It sets `scale_limit` to 1.5. The mask is a registered buffer. The final network layer has zero weights and biases, so initially both scale and shift are zero. This makes the whole coupling layer the identity without preventing later optimization.

The derivation above allows general d; the supplied constructor implements $d = 2$ and accepts only masks $(1, 0)$ and $(0, 1)$. Its network has two hidden layers of width 32 with hyperbolic-tangent activations and four output entries. The method accepts a batch of shape $(B, 2)$ and returns the transformed batch and one log determinant per observation, of shape $(B,)$. Sum over coordinates only. Return the determinant of the applied direction in both branches.

5.2 Worked implementation

```
def transform(self, x, inverse=False):
    fixed = x * self.mask
    s_raw, t_raw = self.net(fixed).chunk(2, dim=-1)
    changed = 1.0 - self.mask
    s = self.scale_limit * torch.tanh(s_raw) * changed
    t = t_raw * changed
    if inverse:
        value = fixed + changed * (x - t) * torch.exp(-s)
        logdet = -s.sum(dim=-1)
    else:
        value = fixed + changed * (x * torch.exp(s) + t)
        logdet = s.sum(dim=-1)
    return value, logdet
```

In the inverse branch, `x` denotes the input to the encoding map, which is the output of the generative map. Its preserved coordinates still supply the same network inputs. A method's argument name need not encode the mathematical direction; the `inverse` flag does.

5.3 What the checks establish

Identity initialization is useful but insufficient: an implementation that always returns its input would pass that check. Nonzero final-layer parameters are needed to test a genuine transformation. The supplied checks examine preserved coordinates, forward–inverse reconstruction, cancellation of the two log determinants at corresponding points, and agreement with an explicit autodifferentiated Jacobian on a small two-dimensional example.

For a single input z, compare the reported forward log determinant with

```
J = torch.autograd.functional.jacobian(
    lambda v: layer.transform(v[None, :])[0][0], z)
sign, reference_logdet = torch.linalg.slogdet(J)
```

Here the two-dimensional Jacobian is used only as a diagnostic. The normal implementation uses the triangular formula rather than building a dense Jacobian for every observation. Double precision makes this small comparison more accurate. A forward–inverse check alone can miss a shared sign error in two wrong log-determinant formulas. The known linear conditioner also gives explicit expected coordinates and determinants. This catches an implementation that consistently drops either the shift or the scale, even if its own inverse and Jacobian still agree.

5.4 Hints and interpretation answers

Start by identifying the preserved coordinates, then evaluate the network once. Mask its outputs before applying either branch. If reconstruction fails, verify that the network receives the preserved part in both directions. If reconstruction succeeds but densities disagree, inspect the determinant sign and the axis of the sum.

Why not invert the neural network? Its inputs are copied unchanged through the layer, so they are already available at the output. Once scale and shift have been recomputed, each changed coordinate has a scalar affine inverse.

Why are zero final-layer parameters acceptable? They initialize the layer as the identity. During the first gradient update, the final layer can receive nonzero gradients from its hidden features. Earlier layers may initially receive zero gradient through zero final weights and begin changing after the final weights move. Initializing every layer to zero would create a different learning problem and is not required here.

Why not calculate a determinant numerically during training? A generic dense determinant costs much more and hides the structural simplification. The explicit small Jacobian is an independent check of the analytic expression.

6 Composition: a RealNVP density model

6.1 Maps, inverses, and order

Compose K generative coupling layers:

$$h_0 = z, \quad h_k = g_k(h_{k-1}), \quad x = h_K = g_K \circ \cdots \circ g_1(z).$$

The chain rule and multiplicativity of determinants give

$$\log |\det Dg(z)| = \sum_{k=1}^{K} \log |\det Dg_k(h_{k-1})|.$$

Encoding applies the last inverse first:

$$z = g_1^{-1} \circ \cdots \circ g_K^{-1}(x).$$

When reading this composition right to left, the first operation on x is g_K^{-1}. In code, iterate through the layers in reverse order and pass `inverse=True` to every layer. Their returned inverse log determinants can simply be accumulated.

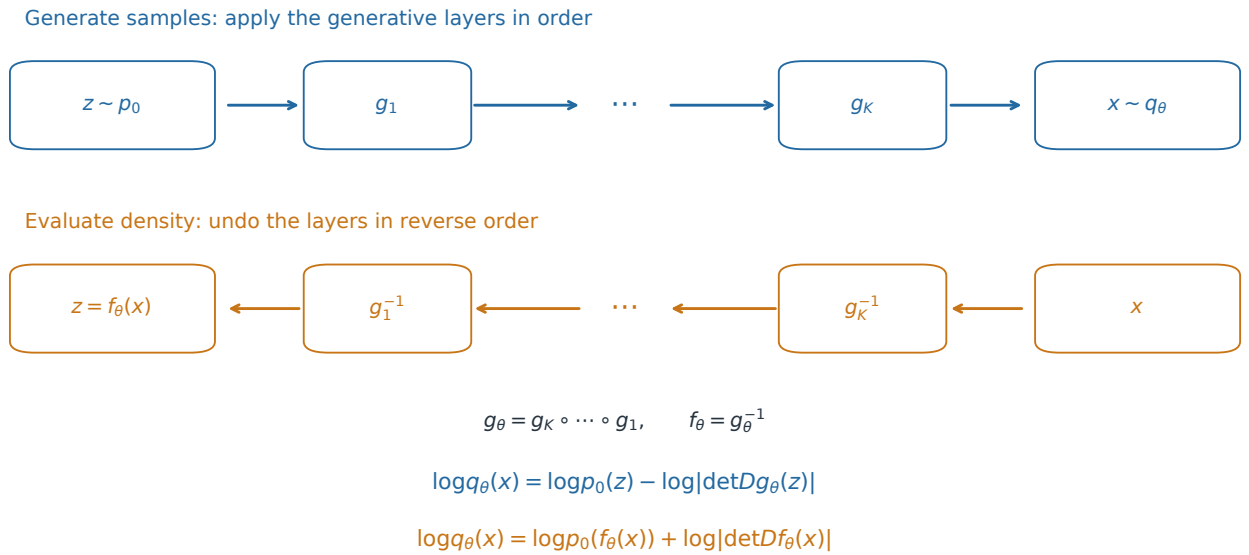


Figure 3: Sampling follows the generative layers from the base to observations. Density evaluation encodes an observation by applying inverse layers in reverse order. Each layer reports the log determinant of the direction actually used.

6.2 Why alternate masks?

With mask $(1, 0)$, the first coordinate is unchanged by a layer. If every layer uses this same mask, the whole composition preserves the first coordinate, so its marginal remains the base marginal. Nonlinear conditional changes in the second coordinate cannot fix a wrong first-coordinate marginal.

The main model uses six layers with alternating masks $(1, 0)$ and $(0, 1)$. Each coordinate can then be transformed using information carried by the other coordinate. The default network has two hidden layers, each of width 32, which keeps the toy problem inexpensive. The number of layers and width are modeling choices; they do not establish that every possible density is represented exactly.

7 Task 3: compose, score, and sample

7.1 What students should do

Complete `RealNVP.transform`, `log_prob`, and `sample(n, generator=None)`. The layers are stored in `self.layers`. The transformation returns $(B, 2)$ values and $(B,)$ log determinants. The supplied base-density helper sums the coordinate log densities along the last axis. It returns $(B,)$ values, not one number for the batch.

7.2 Worked implementation

```
def transform(self, value, inverse=False):
    logdet = value.new_zeros(value.shape[0])
    layers = reversed(self.layers) if inverse else self.layers
    for layer in layers:
        value, contribution = layer.transform(value, inverse=inverse)
        logdet = logdet + contribution
    return value, logdet

def log_prob(self, x):
    z, inverse_logdet = self.transform(x, inverse=True)
    return standard_normal_log_prob(z) + inverse_logdet

@torch.no_grad()
def sample(self, n, generator=None):
    parameter = next(self.parameters())
    z = torch.randn(n, 2, generator=generator, dtype=parameter.dtype,
                    device=parameter.device)
    return self.transform(z, inverse=False)[0]
```

The optional random-number generator must match the model device. The `sample` helper disables gradients for ordinary sampling. Training the variational approximation instead uses `transform(z)` directly.

The standard-normal helper implements

$$\log p_0(z) = -\frac{d}{2}\log(2\pi) - \frac{1}{2}\sum_{j=1}^{d} z_j^2.$$

The inverse output already contains $\log|\det Df(x)|$. Add it when scoring observations. If instead a sample has been generated and the forward log determinant retained, the equivalent score is the base log density *minus* that forward log determinant.

7.3 Checks and model answers

The checks should include a nontrivial composition: all-identity layers cannot expose a reversed-order error. Confirm reconstruction, opposite total log determinants, score agreement along generated points,

the sample shape, and finite values. Gradients of the training objective must reach trainable parameters. Do not put an encoding calculation inside `no_grad` when it is used for training; evaluation and plotting can use that context. The supplied independent composition check uses $g_1(z) = (z_1, 2z_2 + z_1)$ and $g_2(u) = (3u_1 + u_2, u_2)$. The expected output is $(4z_1 + 2z_2, z_1 + 2z_2)$ with log determinant $\log 6$. This catches omitted layers and wrong order even when a round trip is internally consistent.

Why must inverse layer order reverse? For two maps, $g_2(g_1(z)) = x$. Applying g_2^{-1} first gives $g_1(z)$, and then applying g_1^{-1} gives z. Applying inverse maps in the original order generally undoes neither composition.

Are model samples independent? Given fixed fitted parameters, independent standard-normal base draws remain independent under the same deterministic map. The transformed distribution is the learned q_ϕ. There is no accept/reject step or MCMC burn-in in ordinary flow sampling.

8 Likelihood training on noisy two-moon observations

8.1 The objective and its statistical meaning

The supplied data generator creates two semicircles with equal mixing probabilities and adds independent Gaussian coordinate noise of standard deviation 0.08. The labels color plots and are not training inputs. Generate 2,400 training, 600 validation, and 1,000 test observations with independent random seeds. Optimize only with training observations:

$$\widehat{L}(\phi) = -\frac{1}{n_{\text{train}}} \sum_{i=1}^{n_{\text{train}}} \log q_\phi(x_i).$$

The coordinate log densities and log determinants have already been summed inside `log_prob`. The final mean is over observations. PyTorch backpropagates through the explicit inverse and the scale and shift networks; Adam updates their parameters. The fixed main run uses six alternating couplings with hidden width 32, 1,200 steps, minibatches of 256, and learning rate 10^{-3}. Gradient norms are clipped at 5. Full training and validation NLL are evaluated every 100 steps; the best validation state is saved by deep copy and restored before test evaluation. No determinant of a full neural-network Jacobian is needed.

For a population density p_{data} and finite relevant expectations,

$$\mathrm{KL}(p_{\text{data}} \| q_\phi) = \mathbb{E}_{p_{\text{data}}}[\log p_{\text{data}}(X)] - \mathbb{E}_{p_{\text{data}}}[\log q_\phi(X)].$$

The first term does not depend on ϕ, so population maximum likelihood minimizes this forward KL. The empirical loss estimates the cross-entropy term from observations. A finite empirical distribution is a collection of point masses; do not literally call its KL to a continuous density the finite training loss.

```
optimizer.zero_grad()
loss = -flow.log_prob(training_batch).mean()
loss.backward()
optimizer.step()
```

Initialization at the identity starts from the standard-normal density. That is a diagnostic, not a fitted baseline. The Gaussian comparison is fitted independently by maximum likelihood to the training set only. For training observations x_i, its mean and covariance are

$$\widehat{\mu} = \frac{1}{n} \sum_i x_i, \qquad \widehat{\Sigma} = \frac{1}{n} \sum_i (x_i - \widehat{\mu})(x_i - \widehat{\mu})^\top.$$

The MLE covariance uses $1/n$, not the unbiased $1/(n-1)$ factor. The supplied data have a nonsingular fitted covariance, so this baseline does not add a diagonal regularizer. Fit no baseline parameters to validation or test observations.

8.2 Validation and held-out comparison

The training budget and model size are fixed before test evaluation. Validation observations select the checkpoint; the test set is evaluated after that selection. The flow and Gaussian are scored on exactly the same held-out observations in the same original units. The notebook reports negative log likelihood (NLL), in nats per observation, so *lower is better*.

The executed reference run restores step 1,200. Its results are:

Model	Train NLL	Validation NLL	Test NLL	Test SE
Gaussian baseline	1.9011	1.8794	1.9071	0.01779
RealNVP	0.9661	0.9069	0.9743	0.02312

The paired test loss difference, RealNVP minus Gaussian, is -0.9328 with standard error 0.02768. Its negative sign favors the flow. Use the paired losses when discussing uncertainty of this comparison, because both fitted models are evaluated on the same observations.

For pointwise losses $L_i = -\log q(x_i)$, the test standard error is $s_L/\sqrt{n_{\text{test}}}$. For a paired comparison, apply the same formula to $D_i = L_i^{\text{flow}} - L_i^{\text{Gaussian}}$. These conditional standard errors measure variation across new iid test samples with fitted models held fixed. They do not include training randomness, checkpoint selection uncertainty, or variation across new training/validation splits.

The lower flow NLL and its curved geometry favor the flexible model over a single elliptical Gaussian on this problem. A better average score on one finite test set does not certify every tail probability. The numerical values above belong to the executed reference notebook; last digits can vary across software versions. Do not require a universal training-loss threshold or improvement at every optimization step.

If a flow does not outperform the Gaussian in a particular run, inspect the mathematical checks and optimization history before interpreting the result. A correctly evaluated flow density may still underfit, overfit, or be poorly optimized. More training does not repair an incorrect inverse or determinant sign.

8.3 Read the diagnostic panels together

1. **Training and validation loss.** Decreasing loss suggests the optimizer is improving its objective. The plotted training and validation losses are full-data averages evaluated under the same current parameters. Their difference reflects different observations and fit, rather than a comparison of differently timed minibatch values. A validation score slightly better than the training score is possible with finite samples; it does not by itself establish a data leak.
2. **Density contours.** A flow can place density along both curved components. Plotting values on a finite grid does not integrate the full model over $\mathbb{R}^2$. Visual gaps or small grid mass do not prove that the density is unnormalized.
3. **Generated samples.** These show the learned distribution in observation space. Compare geometry, component occupancy, and unusual points with held-out observations. A pleasing scatterplot is not a likelihood check.
4. **Encoded held-out observations.** If $q_\phi = p_{\text{data}}$, applying f_ϕ to observations gives the base distribution. Approximately Gaussian encoded data support the fit, but finite scatterplots and marginal histograms can hide dependence and tail discrepancies. Encoding samples generated by the same model gives the base by construction and principally checks the inverse.

8.4 Worked interpretation and closing prompts

A suitable student explanation identifies the score direction, describes the geometry that the flow learned, and gives a finite-data limitation. For example: "The flow follows the curved components and has a lower held-out NLL in this run. The negative paired loss difference favors the flow, and its standard error measures new-test-sample variation at fixed fitted parameters. More generated samples

reduce integration error under the fitted density; they do not change that density or repair a fitting error."

The encoded test means in the reference run are approximately $(0.0048, -0.0349)$, the marginal variances are $(0.922, 1.042)$, and the covariance is -0.0853. These are useful diagnostics, not exact moment identities for held-out data. A mean near zero alone would not establish a correct joint distribution.

Connect with the preceding seminars. The Laplace approximation uses a Gaussian derived from local posterior curvature. A flow provides a flexible normalized family, but learning its parameters adds an optimization problem. MCMC draws are dependent states designed to preserve a specified target; ordinary fitted-flow draws are independent samples from a learned model. A flow can also be used inside a corrected sampling method, but such a correction is additional machinery and is outside this lesson.

Ask students to answer these three questions at the board:

1. Which direction evaluates the density at an observed x, and why?
2. Which mathematical property permits an ordinary neural network inside an invertible coupling layer?
3. What remains approximate after a flow's Jacobian formula is implemented correctly?

The answers are encoding with the inverse, access to preserved coordinates and a conditional affine inverse, and the learned approximation to the target plus estimation and optimization error.

9 Debugging reference

Observed problem	Check first
Inverse fails after one layer	The network sees only preserved coordinates; the inverse subtracts the shift before multiplying by the inverse scale.
Inverse fails only after composition	Reverse the layer order as well as the direction of each layer.
Inverse works, density fails	Check determinant direction, its sign, and whether the determinant belongs to the actually applied map.
One scalar log determinant for a batch	Sum on the coordinate axis `dim=-1`; retain one value per observation.
No learned change in one coordinate	Inspect masks; identical masks preserve one marginal through the entire composition.
Only identity tests pass	Test after assigning nonzero final-layer parameters; identity initialization hides several implementation errors.
Non-finite loss or gradients	Inspect scale bounds, learning rate, data scale, and the first layer producing non-finite values. Do not silently clip log probabilities to make a failed check pass.
Parameters never change	Verify an optimizer step follows backward, gradients are enabled, and the intended parameters are registered.
CPU/device mismatch	Register masks as buffers; create scalar constants and base samples with matching dtype and device.
Attractive samples, poor test score	Check train-only preprocessing and baseline fitting, inverse scoring, model capacity, and train/validation fit.

A A flow as a Bayesian variational approximation

This appendix is optional and fully supplied; it introduces no fourth implementation task. After completing Tasks 2 and 3, enable it in setup:

```
RUN_APPENDICES = True
```

It reuses the flow and optimizer to approximate a curved, posterior-shaped analytic benchmark, keeping the density calculation separate from the choice of training objective. The target is specified directly and is not the posterior of a fitted data model. Its exact transport, moments, and direct draws provide checks without relying on another approximate inference method.

A.1 A curved, normalized banana target

Let

$$\Theta_1 = Z_1, \qquad \Theta_2 = b(Z_1^2 - 1) + \sigma Z_2, \qquad (Z_1, Z_2) \sim \mathcal{N}(0, I_2),$$

with $b = 0.8$ and $\sigma = 0.45$. Equivalently,

$$\pi(\theta) = \mathcal{N}(\theta_1; 0, 1)\,\mathcal{N}(\theta_2; b(\theta_1^2 - 1), \sigma^2).$$

The exact transport is invertible, with $z_1 = \theta_1$, $z_2 = [\theta_2 - b(\theta_1^2 - 1)]/\sigma$ and constant generative determinant σ. Its log density is

$$\log \pi(\theta) = -\log(2\pi) - \log \sigma - \tfrac{1}{2}\theta_1^2 - \frac{[\theta_2 - b(\theta_1^2 - 1)]^2}{2\sigma^2}.$$

Conditional integration over θ_2 followed by integration over θ_1 shows normalization immediately. The shape is curved, despite both base variables and the conditional observation noise being Gaussian.

A.2 Exact moments and a local Laplace comparison

For a standard-normal variable, $\mathbb{E}Z_1^2 = 1$, $\mathbb{E}Z_1^3 = 0$, and $\mathbb{E}Z_1^4 = 3$. Thus

$$\begin{aligned} \mathbb{E}\Theta_1 &= 0, & \mathbb{E}\Theta_2 &= b(\mathbb{E}Z_1^2 - 1) = 0, \\ \operatorname{Var}\Theta_1 &= 1, & \operatorname{Var}\Theta_2 &= b^2 \operatorname{Var}(Z_1^2) + \sigma^2 = 2b^2 + \sigma^2, \\ \operatorname{Cov}(\Theta_1, \Theta_2) &= b\mathbb{E}[Z_1(Z_1^2 - 1)] = 0. \end{aligned}$$

Consequently

$$\mathbb{E}_\pi \Theta = (0, 0)^\top, \qquad \operatorname{Cov}_\pi \Theta = \operatorname{diag}(1, 1.4825).$$

Zero covariance does not imply independence: the conditional mean of Θ_2 depends on Θ_1^2.

Write $r = \theta_2 - b(\theta_1^2 - 1)$. Apart from a constant, the negative log density is

$$J(\theta) = \tfrac{1}{2}\theta_1^2 + \frac{r^2}{2\sigma^2}.$$

It is minimized uniquely at $\theta_{\mathrm{MAP}} = (0, -b) = (0, -0.8)$, where $r = 0$. At that point $\nabla r = (0, 1)^\top$, so

$$\nabla^2 J(\theta_{\mathrm{MAP}}) = \begin{pmatrix} 1 & 0 \\ 0 & \sigma^{-2} \end{pmatrix}, \qquad q_{\mathrm{L}} = \mathcal{N}\left(\begin{pmatrix} 0 \\ -0.8 \end{pmatrix}, \begin{pmatrix} 1 & 0 \\ 0 & 0.2025 \end{pmatrix} \right).$$

The exact target mean is not its mode. The local Gaussian misses the vertical variation induced by the curved conditional mean. This makes a direct link to the preceding Laplace seminar: positive local curvature yields a valid Gaussian approximation, without guaranteeing global moment accuracy.

A.3 Reverse-KL training

Suppose the target is given by an evaluable kernel $\widetilde{\pi}$ with normalizer Z. A variational objective is

$$\mathrm{KL}(q_\phi \| \pi) = \mathbb{E}_{q_\phi}[\log q_\phi(\Theta) - \log \widetilde{\pi}(\Theta)] + \log Z.$$

The unknown $\log Z$ is constant in ϕ, so it can be omitted while optimizing. Generate $\Theta = g_\phi(Z_0)$ using $Z_0 \sim p_0$. The reparameterized objective is

$$L_{\mathrm{VI}}(\phi) = \mathbb{E}_{Z_0 \sim p_0} \left[\log p_0(Z_0) - \log \left| \det D g_\phi(Z_0) \right| - \log \widetilde{\pi}(g_\phi(Z_0)) \right].$$

The supplied run uses a fresh six-layer, width-32 model, Adam with learning rate 10^{-3}, 1,400 steps, base minibatches of 512, and gradient-norm clipping at 5. Fresh base draws give a stochastic estimate. Backpropagate through the transformed samples, the forward log determinant, and the target evaluation. Do not detach the samples. The base density term is constant with respect to ϕ for a fixed base sample, but retaining it permits an interpretable Monte Carlo objective value.

```
z = torch.randn(batch_size, 2, device=DEVICE, dtype=DTYPE)
theta, forward_logdet = flow.transform(z)
log_q = standard_normal_log_prob(z) - forward_logdet
loss = (log_q - banana_log_prob(theta)).mean()
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

For the fully normalized banana density, this loss estimates $\mathrm{KL}(q_\phi \| \pi)$. If the target function omits a log normalizing constant, it estimates that KL minus $\log Z$, so its numerical value is not itself a KL. Even an estimate of the normalized KL can be slightly negative because a finite Monte Carlo estimate fluctuates; the exact KL is nonnegative.

A.4 How to assess the variational approximation

Compare 20,000 independent samples from each of the trained flow, the exact transport, and the Laplace Gaussian, using new random seeds after training. The supplied table reports means, the second marginal variance, $\Pr(\Theta_2 > 1)$ with its event Monte Carlo standard error, and reverse KL with its integrand Monte Carlo standard error. Do not reuse the last optimization minibatch as an evaluation sample.

The banana's moments above are analytic references. A tail probability estimated from a large exact-transport sample still has Monte Carlo error. Conditional integration also gives the deterministic one-dimensional reference

$$\Pr_\pi(\Theta_2 > c) = \int_{-\infty}^{\infty} \varphi(u) \left[1 - \Phi\left(\frac{c - b(u^2 - 1)}{\sigma} \right) \right] du,$$

where φ and Φ are the standard-normal density and CDF. This integral can be evaluated by numerical quadrature. In the supplied notebook, 320-node Gauss–Hermite quadrature gives $\Pr_\pi(\Theta_2 > 1) \simeq 0.14437009$; 256 nodes give 0.14437012. Their agreement is a useful numerical check, rather than an analytic expression for this probability. State whether a reported reference is analytic, quadrature-based, or sampled.

As an optional algebraic audit, the Laplace approximation in this particular example has

$$\mathrm{KL}(q_{\mathrm{L}} \| \pi) = \mathrm{KL}(\pi \| q_{\mathrm{L}}) = \frac{3b^2}{2\sigma^2} \simeq 4.74074.$$

For the reverse direction, write $\Theta_1 = Z_1$ and $\Theta_2 = -b + \sigma Z_2$ under q_{L}. Then the target residual is $\sigma Z_2 - b Z_1^2$, whose squared expectation is $\sigma^2 + 3b^2$. The log normalizing constants and the first coordinate terms cancel. The forward direction has the same value because under the target $\Theta_2 + b = b Z_1^2 + \sigma Z_2$. Equality of these two KL values is a property of this example, not a general symmetry of KL divergence.

A.5 Reference numerical results and interpretation

The executed run gives the following illustrative evaluation values:

Distribution	Mean 2	Variance 2	$\Pr(\Theta_2 > 1)$	Reverse-KL estimate
Exact target draws	0.0107	1.5204	0.14665	0
Laplace draws	−0.7961	0.2019	0 observed	4.5788
Variational flow	−0.0662	1.1533	0.13595	0.01934

The flow KL estimate has Monte Carlo standard error 0.00192, and its event estimate has standard error 0.00242. The target-reference event estimate has standard error 0.00250. The flow follows the curved mass much better than Laplace, but its second-coordinate variance is still below the analytic value 1.4825. The flow event estimate 0.13595 is also below the quadrature reference 0.14437. Present this remaining error rather than calling the learned target exact or its moments fully recovered.

The Laplace event requires special care. Analytically,

$$\Pr_{q_L}(\Theta_2 > 1) = 1 - \Phi\left(\frac{1+b}{\sigma}\right) = 1 - \Phi(4) \simeq 3.1671 \times 10^{-5}.$$

With 20,000 draws the expected event count is only 0.6334. Zero observed events, and hence a zero plug-in Bernoulli standard error in the displayed table, are quite plausible. Neither implies that the event probability or its estimation uncertainty is zero. The Laplace KL estimate 4.5788 has Monte Carlo standard error 0.1042; compare it with the analytic value 4.7407, allowing for sampling variation.

Interpretation answer. The target mean averages the entire curved mass and need not equal the density mode. The trained flow improves the geometry and many summaries relative to Laplace, while retaining visible approximation error. Its model density is exactly evaluable even when its moments or tails disagree with the target. More draws reduce numerical error under the fitted flow; they do not eliminate this discrepancy. With an unknown target normalizer, the same VI gradients remain available, but the absolute KL value cannot be recovered from the kernel objective alone. Exact target moments and independent reference draws would usually be unavailable too.

B Support, expressiveness, and the direction of KL

B.1 Support constraints

A finite composition of the stated coupling bijections maps $\mathbb{R}^d$ onto $\mathbb{R}^d$. With a positive standard-normal base, the model density is positive throughout this space. It cannot assign exactly zero density to an open gap or produce a distribution exactly supported on a lower-dimensional curve. Noisy moons provide an ordinary continuous example; an ideal noiseless one-dimensional curve in $\mathbb{R}^2$ has no two-dimensional Lebesgue density.

A connected full-support model can nevertheless place very little probability between separated high-density regions and can represent multiple density modes. "A bijection cannot produce a multimodal density from a Gaussian" is false: the spatially varying Jacobian changes density heights. What remains constrained here is the exact support, as well as the finite model's capacity. Targets restricted to a positive domain or a simplex require an appropriate additional support transformation and its Jacobian.

B.2 Computational and modeling tradeoffs

The coupling structure makes both the inverse and determinant inexpensive. A single layer leaves some coordinates unchanged, so composition and changed masks are essential. Additional layers and network width increase flexibility and computational cost. Bounded scales help numerics while also restricting what a fixed number of layers can express. None of these choices removes the need for held-out evaluation or a target-specific adequacy check.

Density values depend on units: rescaling coordinates changes a log density by the corresponding log-Jacobian term. Compare models on the same transformed data, or consistently convert all scores back

to original units. Differential log densities can be positive, and a negative average negative log likelihood is not by itself an error.

B.3 Likelihood and variational objectives ask different questions

Training setting	Objective and source of samples	Main implication
Density estimation	$-\mathbb{E}_{p_{\text{data}}} \log q_\phi$; samples come from observations.	Population objective differs from $\text{KL}(p_{\text{data}} \| q_\phi)$ by a constant.
Variational inference	$\mathbb{E}_{q_\phi}[\log q_\phi - \log \widetilde{\pi}]$; samples come from the flow.	Objective differs from $\text{KL}(q_\phi \| \pi)$ by a constant independent of ϕ.

Forward KL penalizes assigning too little density where the target supplies probability mass. Reverse KL averages over where the current approximation places its own mass; missed regions may contribute little to its sampled objective. These observations explain familiar tendencies, but do not justify blanket promises that likelihood covers every mode or reverse KL always selects one mode. Model family, support, target geometry, data coverage, and optimization all matter. For an exactly representable target, a global minimum of either KL is obtained at $q_\phi = \pi$ when the divergences are well-defined.

Final oral prompt. Can more samples from a fitted flow repair a wrong variational approximation? More samples reduce Monte Carlo error for expectations under that fitted q_ϕ, under the usual moment assumptions. They do not change q_ϕ or eliminate its discrepancy from the target. That is the same distinction between numerical uncertainty and inferential approximation emphasized in the sampling and Laplace seminars.

References

[1] Laurent Dinh, Jascha Sohl-Dickstein, and Samy Bengio. *Density estimation using Real NVP*. ICLR, 2017. arXiv:1605.08803.

[2] Danilo Jimenez Rezende and Shakir Mohamed. *Variational inference with normalizing flows*. ICML, 2015. arXiv:1505.05770.